



Institut et hôpital neurologiques de Montréal
Montreal Neurological Institute and Hospital

MCIN MCGILL CENTRE
for INTEGRATIVE
NEUROSCIENCE




CENTRE
LUDMER
NEUROINFORMATIQUE & SANTÉ MENTALE | NEUROINFORMATICS & MENTAL HEALTH

Intro to Deep Learning for NeuroImaging

Andrew Doyle

McGill Centre for Integrative Neuroscience

[@crocodoyle](#)

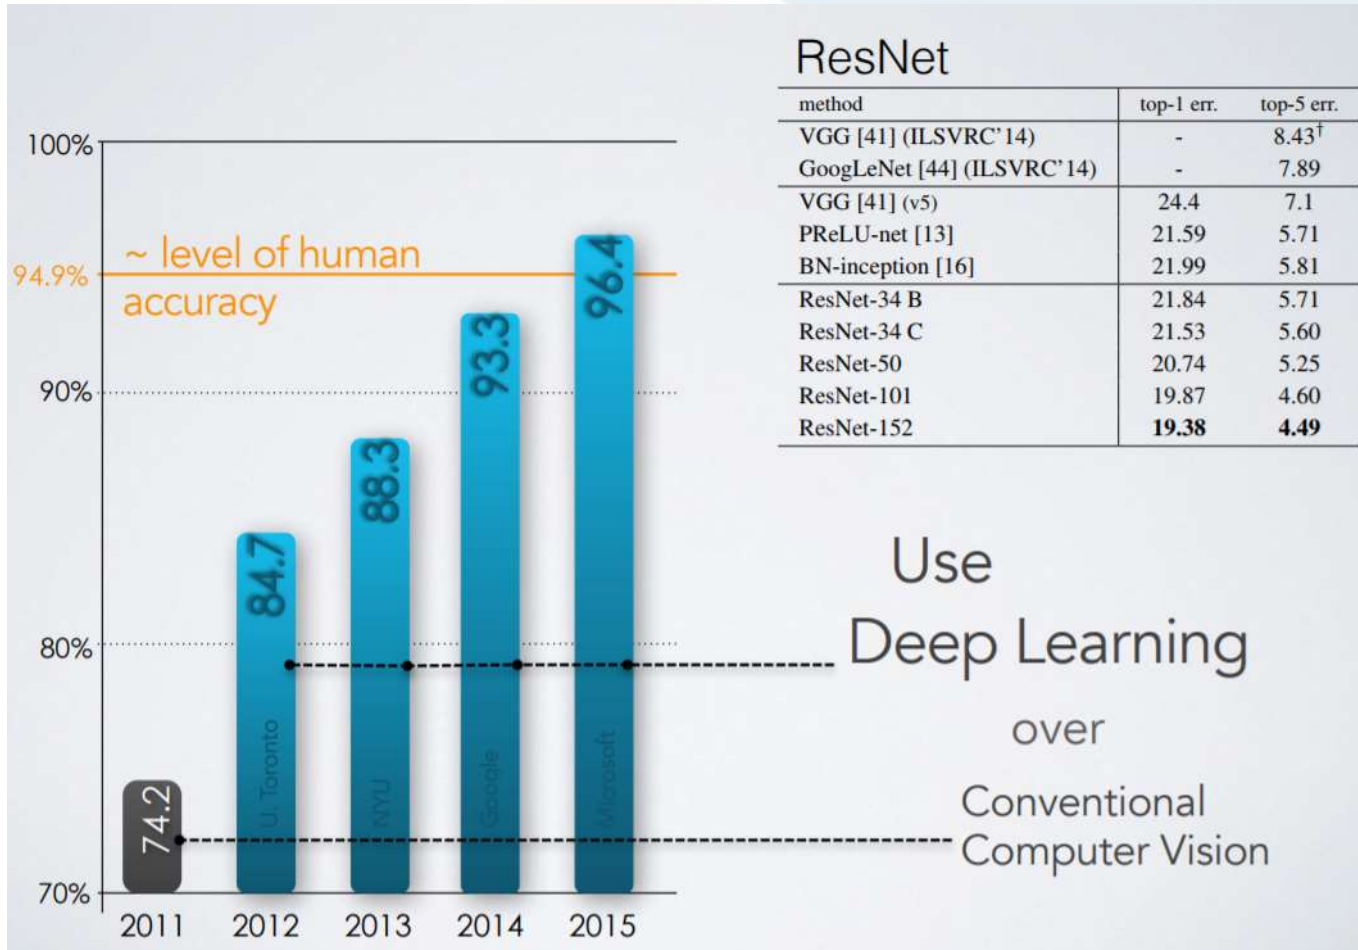


Outline


1. GET EXCITED
2. Artificial Neural Networks
3. Backpropagation
4. Convolutional Neural Networks
5. Neuroimaging Applications




ImageNet-1000 Results



Generative Models




BrainBrush



Deep Blood by Team BloodArt

Gatys, Leon A., Alexander S. Ecker, and Matthias Bethge. "Image style transfer using convolutional neural networks." *Computer Vision and Pattern Recognition (CVPR), 2016 IEEE Conference on*. IEEE, 2016.

Generative Models

Text description	This bird is blue with white and has a very short beak	This bird has wings that are brown and has a yellow belly	A white bird with a black crown and yellow beak	This bird is white, black, and brown in color, with a brown beak	The bird has small beak, with reddish brown crown and gray belly	This is a small, black bird with a white breast and white on the wingbars.	This bird is white black and yellow in color, with a short black beak
Stage-I images							
Stage-II images							

StackGAN

Generative Models



CycleGAN

Generative Models

MR




CT



CycleGAN


Wolterink, Jelmer M., et al. "Deep MR to CT synthesis using unpaired data." *International Workshop on Simulation and Synthesis in Medical Imaging*. Springer, Cham, 2017.

Synthetic Images



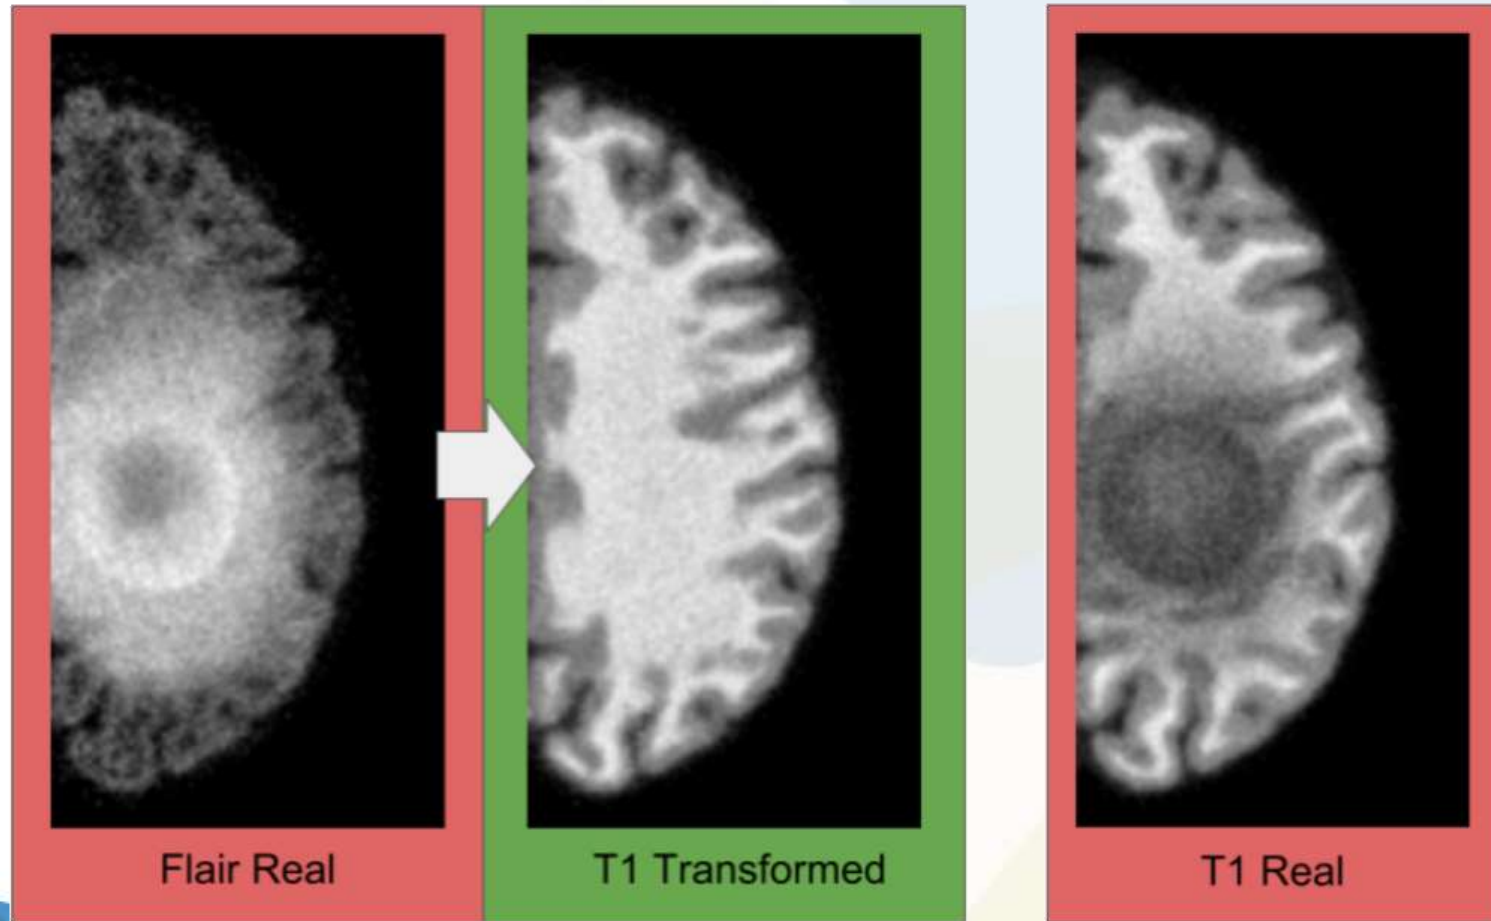
Real

Synthesized



FLAIR

Synthetic Images



Reinforcement Learning





Reinforcement Learning OpenAI





Reinforcement Learning OpenAI



mini-map

Montreal AI Companies



ARCTIC FOX AI



facebook research

Deep Learning


For Deep Learning, you need:

1. Artificial Neural Network
2. Loss
3. Optimizer
4. Data

Artificial Neurons

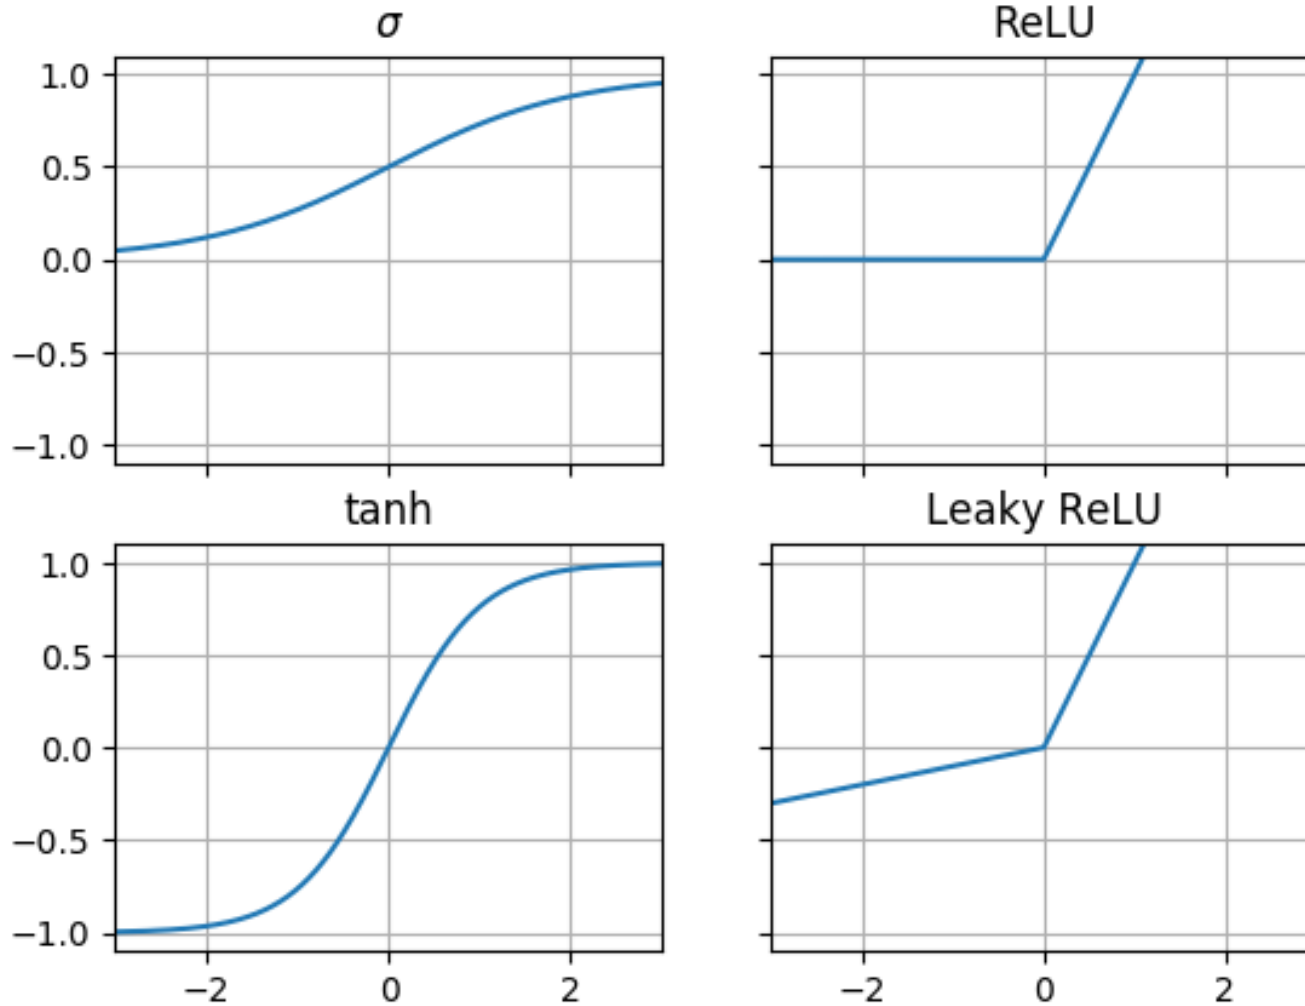


Artificial Neurons




$$o = f(x) = f(\mathbf{w}^T \mathbf{i} + \mathbf{b})$$

Artificial Neurons




Artificial Neurons


Logistic Regression





$$o = \sigma(x) = \sigma(\mathbf{w}^T \mathbf{i} + b)$$





Neural Networks



Neural Networks




Neural Networks



Neural Networks



Neural Networks



$$f(y) = \sigma(w_{y,h_1}f(h_1) + w_{y,h_2}f(h_2) + b_y)$$

$$f(x_2) = \sigma(i_2w_{x_2,i_2} + b_{x_2}) = \sigma(w_{y,h_1}\sigma(w_{h_1,x_1}\sigma(i_1w_{x_1,i_1} + b_{x_1}) + w_{h_1,x_2}\sigma(i_2w_{x_2,i_2} + b_{x_2}) + b_{h_1}))$$

$$f(h_2) = \sigma(w_{h_2,x_1}f(x_1) + w_{h_2,x_2}f(x_2) + b_{h_2}) = \sigma(w_{h_2,x_1}\sigma(i_1w_{x_1,i_1} + b_{x_1}) + w_{h_2,x_2}\sigma(i_2w_{x_2,i_2} + b_{x_2}) + b_{h_2})$$

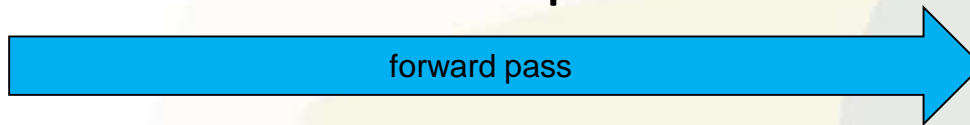
17 parameters $\theta = \{w, b\}$

Backpropagation

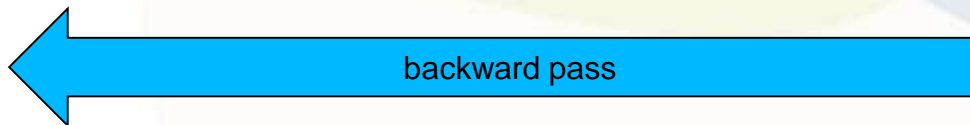
1. Random θ initialization

Iterate:


1. Forward - compute loss



2. Backward - update parameters



Backpropagation



forward pass

$$L(o, \hat{y}) = \frac{1}{2} \sum (o - \hat{y})^2$$



backward pass

XOR

i_1	i_2	o
0	0	0
0	1	1
1	0	1
1	1	0

$$\nabla_{\theta} L(o, \hat{y}) = \left[\frac{\partial L}{\partial w_{x_1, i_1}}, \frac{\partial L}{\partial b_{x_1}}, \frac{\partial L}{\partial w_{x_2, i_2}}, \frac{\partial L}{\partial b_{x_2}}, \dots, \frac{\partial L}{\partial w_{y, h_2}} \right]^T$$

Backpropagation



Gradients in blue

Backpropagation

Initialize w

forward pass

backward pass

J


$\frac{\partial L}{\partial w}$

$$w' = w - \alpha \frac{\partial L}{\partial w}$$

learning rate

w

Backpropagation




$$\frac{\partial L}{\partial w_{y,h_1}} = \frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_{y,h_1}}$$

...

$$= \sum -\sigma(\hat{y})(1 - \sigma(\hat{y})) f(h_1)$$


Backpropagation



$$\frac{\partial L}{\partial w_{h_1, x_1}} = \frac{\partial L}{\partial y} * \frac{\partial y}{\partial h_1} * \frac{\partial h_1}{\partial w_{h_1, x_1}}$$

$$\frac{\partial L}{\partial w_{h_2, x_2}} = \frac{\partial L}{\partial y} * \frac{\partial y}{\partial h_2} * \frac{\partial h_2}{\partial w_{h_2, x_2}}$$

Backpropagation



$$\frac{\partial L}{\partial w_{x_1, i_1}} = \frac{\partial L}{\partial y} * \frac{\partial y}{\partial h_1} * \frac{\partial h_1}{\partial x_1} * \frac{\partial x_1}{\partial w_{x_1, i_1}} + \frac{\partial L}{\partial y} * \frac{\partial y}{\partial h_2} * \frac{\partial h_2}{\partial x_1} * \frac{\partial x_1}{\partial w_{x_1, i_1}}$$



TensorFlow

PYTORCH

Optimizers

1. Gradient Descent

$$w' = w - \alpha \frac{\partial J}{\partial w}$$

2. Stochastic Gradient Descent

approx. $\frac{\partial J}{\partial w}$ in batches

3. Momentum

$$w' = w + v \quad v = \gamma v + \alpha \frac{\partial J}{\partial w}$$

4. Adagrad/adadelata

param-wise decaying learning rate

5. RMSprop

avg. gradients

6. Adam

RMSprop + momentum


Data Manifold

Data distribution:

- Class 1
- Class 2

X-Y grid:

- Param (θ) space




Data Manifold

Data distribution:

- Class 1
- Class 2

X-Y grid:

- Param (θ) space



Convolutional Neural Networks

$$f(t) * g(t) = \int_{\tau=-\infty}^{\infty} f(\tau) \cdot g(t - \tau) d\tau$$


Convolutional Neural Networks

Convolution of a block with a block

$g(t)$

$f(t)$

$f(t) * g(t)$




Convolutional Neural Networks

Convolution of a triangle with a block


$g(t)$

$f(t)$


$f(t) * g(t)$



Convolutional Neural Networks



Convolutional Neural Networks




Convolutional Neural Networks

CNN/convnet neurons:


1. Have receptive fields
2. Share weights
 - Vastly reduces parameters
 - Translational equivariance

Convolutional Neural Networks




Convolution

Convolutional Neural Networks




Convolution

Convolutional Neural Networks




Convolution

Convolutional Neural Networks




Convolution

Convolutional Neural Networks



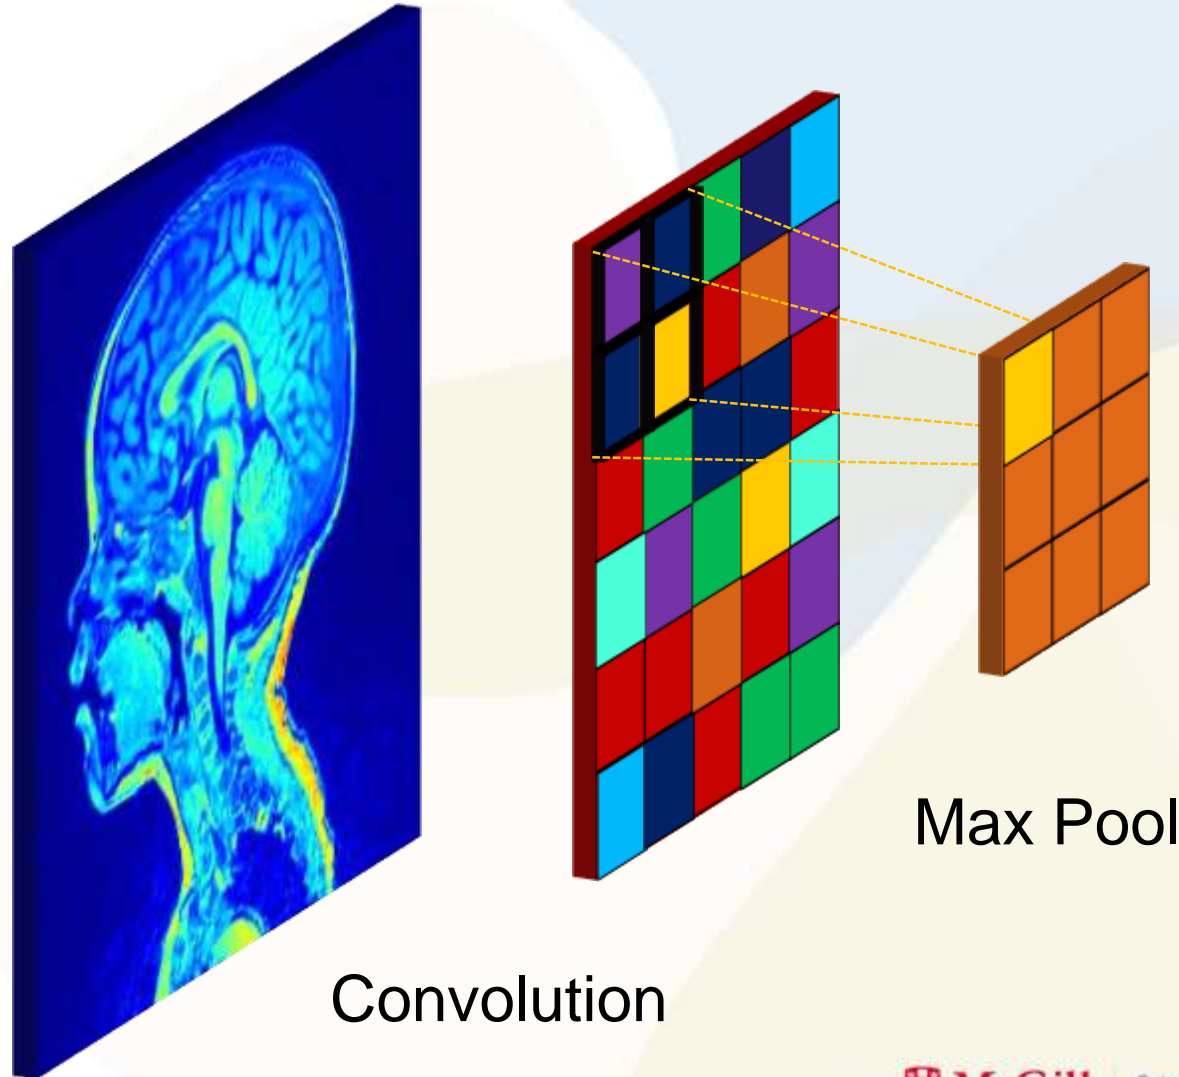
Convolution

Convolutional Neural Networks

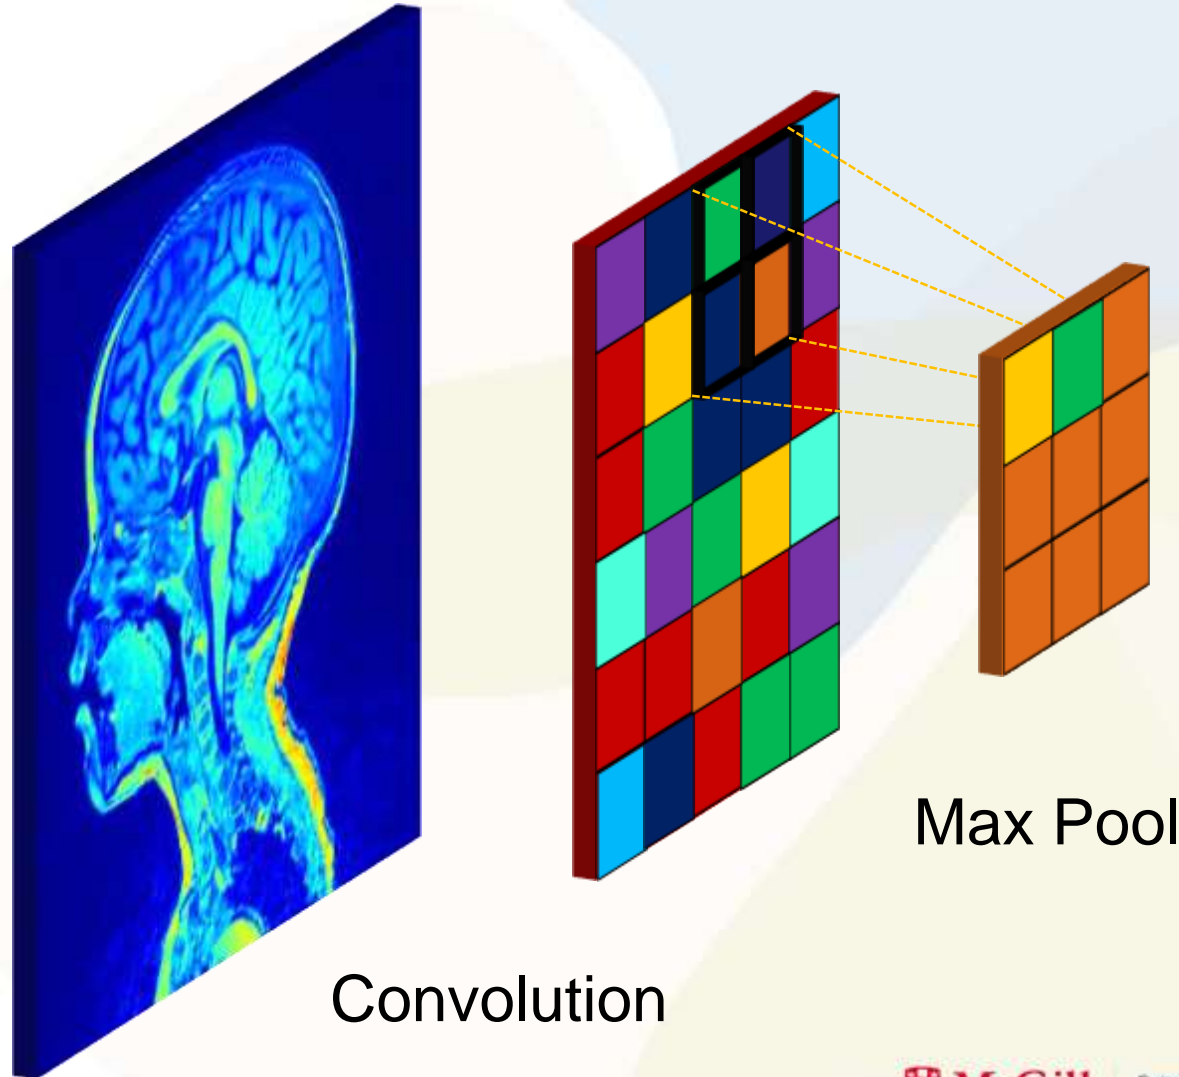


Convolution


Convolutional Neural Networks




Convolutional Neural Networks



Convolutional Neural Networks




Convolution



Max Pool


Convolutional Neural Networks




AlexNet trained using **Dropout**

90% parameters

Convolutional Neural Networks




Dropout




$$P(\text{dropout}) = 0.5$$

Dropout




$$P(\text{dropout}) = 0.5$$

Dropout



$$P(\text{dropout}) = 0.5$$

Batch Normalization




- Subtract mean, divide by standard deviation

Batch Normalization

- Whitens activations
- Speeds training
- Injects noise
- Not good for small batch sizes


Group Normalization



- N : training examples
- C : channels
- H, W : spatial dimensions

Convolutional Neural Networks

- 3x3 convolutions




VGG16

ResNet

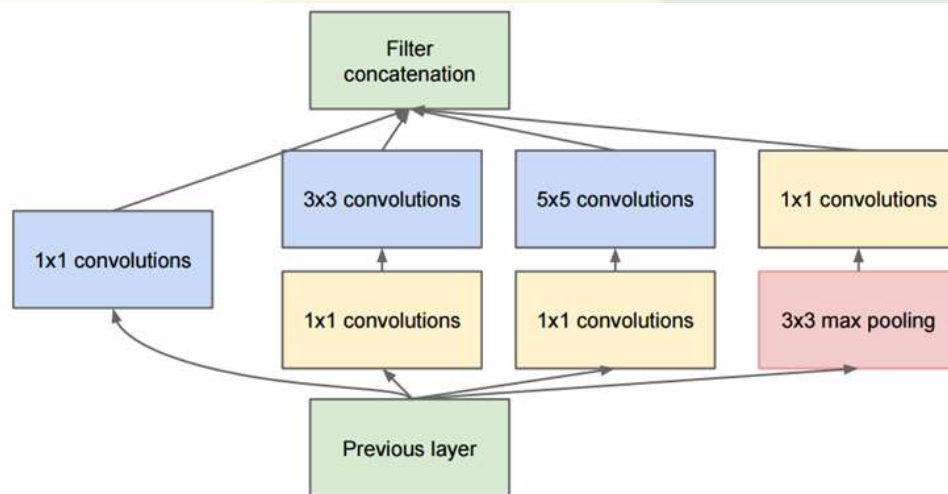
152 convolutional layers
Skip (residual) connections

34-layer residual

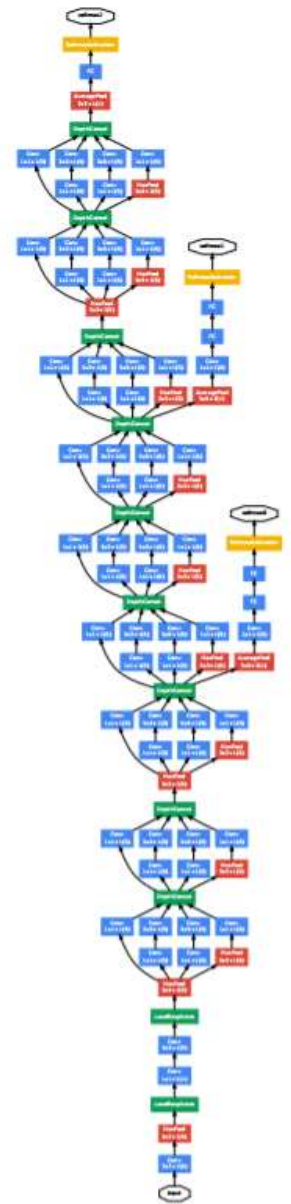


GoogLeNet


1. Deep Supervision helps training
2. 1x1 convolutions can replace fully-connected layers




Full Inception module



DenseNet



DenseNet



- Densely-connected blocks & transition layers
- Far more parameter-efficient
- Doesn't need fancy optimizers

Challenges


1. Data quantity
2. Data size
3. Data quality
4. Data variability
5. Unexpected pathology

Start here



Beyond Linear Decoding
Introduction to Deep Learning

bit.ly/ohbmdl



Keras

keras.io

deeplearningbook.org

Thanks!

